

```

\ COPYRIGHT 1991 JULIAN V. NOBLE
TASK INTEGRAL
FIND CP@L 0 = ?( FLOAD COMPLEX )
\ define data-type tokens if not already
FIND REAL*4 0 = ?(((
    0 CONSTANT REAL*4
    1 CONSTANT REAL*8
    2 CONSTANT COMPLEX
    3 CONSTANT DCOMPLEX )))

FIND 1ARRAY 0 = ?( FLOAD MATRIX.HSF )
\ function usage
: USE( [COMPILE] ' CFA ; IMMEDIATE

\ BEHEADING starts here
0 VAR N

: inc.N N 1+ IS N ;
: dec.N N 2- IS N ;

0 VAR type

\ define "stack"
    20 LONG REAL*8 1ARRAY X{
    20 LONG REAL*4 1ARRAY E{
    20 LONG DCOMPLEX 1ARRAY F{
    20 LONG DCOMPLEX 1ARRAY I{

2 DCOMPLEX SCALARS old.I final.I
:)integral ( n-- )\ trapezoidal rule
    X{ OVER } G@L
    X{ OVER 1- } G@L
    F- F2/
    F{ OVER } G@L
    F{ OVER 1- } G@L
    type 2 AND
    IF X+ FROT X*F
    ELSE F+ F* THEN
    I{ SWAP 1- } G!L ;

0 VAR f.name
: f(x) f.name EXECUTE ;

```

```

: INITIALIZE
    IS type \ store type
    type F{ I
    type I{ I\ set types for function
    type 'old.I I
    type 'final.I I \ and integral(s)
    type 1 AND X{ I
        \ set type for indep. var.
    E{ 0 } G!L \ store error
    X{ 1 } G!L \ store B
    X{ 0 } G!L \ store A
    IS f.name \ ! cfa of f(x)
    X{ 0 } G@L f(x) F{ 0 } G!L
    X{ 1 } G@L f(x) F{ 1 } G!L
    1 IS N
    N)integral
    type 2 AND IF F=0 THEN
    F=0 final.I G!L
    FINIT ;

: E/2 E{ N 1- } G@L F2/ E{ N 1- } G!L ;

: }move.down ( adr n-- )
    } #BYTES >R (-- seg off)
    DDUP R@ +
    (-- s.seg s.off d.seg d.off)
    R> CMOVE! ;

: MOVE.DOWN
    E{ N 1- }move.down
    X{ N }move.down
    F{ N }move.down ;

: new.X ( 87:--x' )
    X{ N } G@L X{ N 1- } G@L
    F+ F2/ FDUP X{ N } G!L ;

\ cont'd. ...

```

```

\ INTEGRAL cont'd
\ debugging code
: GF. 1 > IF FSWAP E. THEN E. ;
: F@. DUP > R G@L R > GF. ;
: .STACKS CR ." N"
      8 CTAB ." X"
      19 CTAB ." Re[F(X)]"
      31 CTAB ." Im[F(X)]"
      45 CTAB ." Re[I]"
      57 CTAB ." Im[I]"
      71 CTAB ." E"
      N 2 + 0 DO CR I.
          3 CTAB X{1} F@.
          16 CTAB F{1} F@.
          42 CTAB I{1} F@.
          65 CTAB E{1} F@.
      LOOP
      CR 5 SPACES ." old.I =" old.I F@.
      5 SPACES ." final.I =" final.I F@. CR ;

```

```

CASE: <DEBUG> NEXT .STACKS ;CASE
0 VAR (DEBUG)
: DEBUG-ON 1 IS (DEBUG) 5 #PLACES ! ;
: DEBUG-OFF 0 IS (DEBUG) 7 #PLACES ! ;
: DEBUG (DEBUG) <DEBUG> ;

```

```

: SUBDMIDE
  N 19 > ABORT" Too many subdivisions!"
  E/2 MOVE.DOWN
  I{N 1-} DROP old.I #BYTES CMOVE
  new.X f(x) F{N} GIL
  N)integral N 1 + )integral ;

```

```

: CONVERGED? ( 87: -- I[N] + I'[N-1] - I[N-1] : -- )
  I{N} G@L I{N 1-} G@L old.I G@L
  type 2 AND
  IF CP- CP+ CPDUP CPABS
  ELSE F- F+ FDUP FABS
  THEN
  E{N 1-} G@L F2* F< ;

```

```

CASE: g*f CP*F F* ;CASE
4 S->F 3 S->F F/ FCONSTANT F=4/3

```

```

: INTERPOLATE ( 87: I[N] + I'[N-1] - I[N-1] : -- )
  F=4/3 type 2/ g*f
  old.I G@L final.I G@L
  type 2 AND
  IF CP+ CP+
  ELSE F+ F+ THEN
  final.I GIL ;
\ BEHEADING ends here

```

```

: )INTEGRAL ( 87: A B ERR -- I[A,B] )
  INITIALIZE
  BEGIN N 0 >
  WHILE SUBDMIDE DEBUG
    CONVERGED? Inc.N
    IF INTERPOLATE dec.N
    ELSE type 2 AND IF FDROP
    THEN FDROP
  THEN
    REPEAT final.I G@L ;
  BEHEAD" N INTERPOLATE \ optional
  \ USE( F.name % A % B % E type )INTEGRAL

```

The nonrecursive program obviously requires *much* more code than the recursive version. This is the chief disadvantage of a nonrecursive method¹⁶.

16. The memory usage is about the same: the recursive method pushes limits, etc. onto the fstack.

§§5-4 Example of)INTEGRAL IN USE

The debugging code ("DEBUG-ON") lets us track the execution of the program by exhibiting the simulated stacks. Here is an

example, $\int_1^2 dx \sqrt{x}$:

USE(FSQRT % 1. % 2. % 1.E-3 REAL*4)INTEGRAL E

```
N      X      F      I      E
0 1.0000E+00 1.0000E+00 5.5618E-01 5.0000E-04
1 1.5000E+00 1.2247E+00 6.5973E-01 5.0000E-04
2 2.0000E+00 1.4142E+00 1.4983E-01 1.2500E-04
old.I = 1.2071E+00 final.I = 0.0000E+00
```

```
0 1.0000E+00 1.0000E+00 5.5618E-01 5.0000E-04
1 1.5000E+00 1.2247E+00 3.1845E-01 2.5000E-04
2 1.7500E+00 1.3228E+00 3.4213E-01 2.5000E-04
3 2.0000E+00 1.4142E+00 1.7396E-01 1.2500E-04
old.I = 6.5973E-01 final.I = 0.0000E+00
```

```
0 1.0000E+00 1.0000E+00 5.5618E-01 5.0000E-04
1 1.5000E+00 1.2247E+00 3.1845E-01 2.5000E-04
2 1.7500E+00 1.3228E+00 1.6826E-01 1.2500E-04
3 1.8750E+00 1.3693E+00 1.7396E-01 1.2500E-04
4 2.0000E+00 1.4142E+00 0.0000E+00 0.0000E+00
old.I = 3.4213E-01 final.I = 0.0000E+00
```

```
0 1.0000E+00 1.0000E+00 5.5618E-01 5.0000E-04
1 1.5000E+00 1.2247E+00 1.5621E-01 1.2500E-04
2 1.6250E+00 1.2747E+00 1.6235E-01 1.2500E-04
3 1.7500E+00 1.3228E+00 1.7396E-01 1.2500E-04
old.I = 3.1845E-01 final.I = 3.4226E-01
```

```
N      X      F      I      E
0 1.0000E+00 1.0000E+00 2.6475E-01 2.5000E-04
1 1.2500E+00 1.1180E+00 2.9284E-01 2.5000E-04
2 1.5000E+00 1.2247E+00 1.6235E-01 1.2500E-04
old.I = 5.5618E-01 final.I = 6.6087E-01
```

```
0 1.0000E+00 1.0000E+00 2.6475E-01 2.5000E-04
1 1.2500E+00 1.1180E+00 1.4316E-01 1.2500E-04
2 1.3750E+00 1.1726E+00 1.4983E-01 1.2500E-04
3 1.5000E+00 1.2247E+00 1.7396E-01 1.2500E-04
old.I = 2.9284E-01 final.I = 6.6087E-01
```

```
0 1.0000E+00 1.0000E+00 1.2879E-01 1.2500E-04
1 1.1250E+00 1.0606E+00 1.3616E-01 1.2500E-04
2 1.2500E+00 1.1180E+00 1.4983E-01 1.2500E-04
old.I = 2.6475E-01 final.I = 9.5392E-01
1.2189E+00 ok
```

$$\int_1^2 dx \sqrt{x} = \frac{2}{3} (2^{3/2} - 1)$$

Notice that, although \sqrt{x} is perfectly finite at $x = 0$, its first derivative is not. This is not a problem in the above case, because the lower limit is 1.0.

It is an instructive exercise to run the above example with the limits (0.0, 1.0). The adaptive routine spends many iterations approaching $x = 0$ (25 in the range [0., 0.0625] vs. 25 in the range [0.0625, 1.0]). This is a concrete example of how an adaptive routine will unerringly locate the (integrable) singularities of a function by spending lots of time near them. The best answer to this problem is to separate out the bad parts of a function by hand, if possible, and integrate them by some other algorithm that takes the singularities into account. By the same token, one should always integrate *up to*, but not *through*, a discontinuity in $f(x)$.

§§6 Adaptive integration in the Argand plane

We often want to evaluate the complex integral

$$I = \oint_{\Gamma} f(z) dz \quad (16)$$

where Γ is a contour (simple, closed, piecewise-continuous curve) in the complex z -plane, and $f(z)$ is an **analytic**¹⁷ function of z .

The easiest way to evaluate 16 is to parameterize z as a function of a real variable t ; as t runs from A to B , $z(t)$ traces out the contour. For example, the parameterization

$$z(t) = z_0 + R \cos(t) + iR \sin(t), \quad 0 \leq t \leq 2\pi \quad (17)$$

traces out a (closed) circle of radius R centered at $z = z_0$.

We assume that the derivative $\dot{z}(t) \equiv \frac{dz}{dt}$ can be defined; then the integral 16 can be re-written as one over a real interval, with a complex integrand:

$$I = \int_A^B \dot{z}(t) f(z(t)) dt \quad (18)$$

Now our previously defined adaptive function **INTEGRAL** can be applied directly, with **F.name** the name of a *complex* function

$$g(t) = \dot{z}(t) f(z(t)), \quad (19)$$

of the *real* variable t .

Here is an example of complex integration: we integrate the function $f(z) = e^{1/z}$ around the unit circle in the counter-clockwise (positive) direction.

7. "Analytic" means the ordinary derivative $df(z)/dz$ exists. Consult any good text on the theory of functions of a complex variable.

The calculus of residues (Cauchy's theorem) gives

$$\oint_{|z|=1} dz e^{1/z} = 2\pi i \quad (20)$$

We parameterize the unit circle as $z(t) = \cos(2\pi t) + i \sin(2\pi t)$ hence $\dot{z}(t) = 2\pi i z(t)$, and we might as well evaluate

$$\int_0^1 dt z(t) e^{1/z(t)} \equiv 1. \quad (21)$$

For reasons of space, we exhibit only the first and last few iterations

```

FIND FSINCOS 0 = ?( FLOAD TRIG )
: Z(T) F=PI F* F2* FSINCOS ;      ( 87: t -- z )
: XEXP FSINCOS FROT FEXP X*F ;
  ( 87: x y -- e^x cos[y] e^x sin[y] )
: G(T) Z(T) XDUP 1/X XEXP X* ;
DEBUG-ON
USE( G(T) % 0 % 1 % 1.E-2 COMPLEX ) INTEGRAL X.

```

N	X	F	I	E
0	0.0000	2.7182	0.0000	.58760 0.0000 .0049999
1	.50000	-.36787	-0.0000	.58760 0.0000 .0049999
2	1.0000	2.7182	0.0000	.00000 .00000 .0000000

old.I = 2.7182 0.0000 final.I = 0.0000 0.0000

N	X	F	I	E
0	0.0000	2.7182	0.0000	.58760 0.0000 .0049999
1	.50000	-.36787	-0.0000	.059198 -.067537 .0024999
2	.75000	.84147	-.54030	.44496 -.067537 .0024999
3	1.0000	2.7182	0.0000	.00000 .00000 .0000000

old.I = .58760 0.0000 final.I = 0.0000 0.0000

N	X	F	I	E
0	0.0000	2.7182	0.0000	.58760 0.0000 .0049999
1	.50000	-.36787	-0.0000	.059198 -.067537 .0024999
2	.75000	.84147	-.54030	.17896 -.043682 .0012499
3	.87500	2.0219	-.15862	.29626 -.008914 .0012499
4	1.0000	2.7182	0.0000	.00000 .00000 .0000000

old.I = .44496 -.067537 final.I = 0.0000 0.0000

N	X	F	I	E
0	0.0000	2.7182	0.0000	.58760 0.0000 .0049999
1	.50000	-.36787	-0.0000	.059198 -.067537 .0024999
2	.75000	.84147	-.54030	.17896 -.043682 .0012499
3	.87500	2.0219	-.15862	.14190 -.005745 .00082499
4	.93750	2.5189	-.025229	.16366 -.000788 .00082499
5	1.0000	2.7182	0.0000	.000000 .000000 .0000000

old.I = .29626 -.0089138 final.I = 0.0000 0.0000

N	X	F	I	E
0	0.0000	2.7182	0.0000	.58760 0.0000 .0049999
1	.50000	-.36787	-0.0000	.059198 -.067537 .0024999
2	.75000	.84147	-.54030	.17896 -.043682 .0012499
3	.87500	2.0219	-.15862	.14190 -.005745 .00082499
4	.93750	2.5189	-.025229	.08102 -.0004467 .00031249
5	.96875	2.6665	-.0033577	.08414 -.0000525 .00031249
6	1.0000	2.7182	0.0000	.000000 .000000 .00000000

old.I = .16366 -.00078842 final.I = 0.0000 0.0000

N	X	F	I	E
0	0.0000	2.7182	0.0000	.58760 0.0000 .0049999
1	.50000	-.36787	-0.0000	.059198 -.067537 .0024999
2	.75000	.84147	-.54030	.17896 -.043682 .0012499
3	.87500	2.0219	-.15862	.14190 -.0057453 .00082499
4	.93750	2.5189	-.025229	.08102 -.0004467 .00031249
5	.96875	2.6665	-.0033577	.04197 -.0000296 .00015624
6	.98437	2.7052	-.000426	.042371 -.0000033 .00015624
7	1.0000	2.7182	0.0000	0.0000 0.0000 0.0000

old.I = .084137 -.000052465 final.I = 0.0000 0.0000

N	X	F	I	E
0	0.0000	2.7182	0.0000	.58760 0.0000 .0049999
1	.50000	-.36787	-0.0000	.059198 -.067537 .0024999
2	.75000	.84147	-.54030	.17896 -.043682 .0012499
3	.87500	2.0219	-.15862	.14190 -.005745 .00082499
4	.93750	2.5189	-.025229	.040020 -.0002833 .00015624
5	.95312	2.6036	-.011038	.041173 -.0001125 .00015624
6	.96875	2.6665	-.003358	.042371 -.0000033 .00015624
7	1.0000	2.7182	0.0000	.000000 .000000 .0000000

old.I = .081022 -.00044667 final.I = .084404 -.0000263

N	X	F	I	E
0	0.0000	2.7182	0.0000	.29826 .0089138 .0012499
1	.12500	2.0219	.15862	.10753 .016479 .00082499
2	.18750	1.4180	.36873	.039711 .013045 .00031249
3	.21875	1.1224	.49820	.030886 .015726 .00031249
4	.25000	.84147	.54030	.000053670 .0079678 .00015624
old.I = .070842 .028407 final.I = .51343 -.050786N				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.29826 .0089138 .0012499
1	.12500	2.0219	.15862	.058518 .0085556 .00031249
2	.15625	1.7232	.28094	.049099 .0086386 .00031249
3	.18750	1.4180	.36873	.030886 .015726 .00031249
old.I = .10753 .016479 final.I = .58374 -.021892				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.16386 .00078842 .00062499
1	.062500	2.5189	.025229	.14180 .0057453 .00062499
2	.12500	2.0219	.15862	.049099 .0086386 .00031249
old.I = .29826 .0089138 final.I = .69139 -.0055268				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.16386 .00078842 .00062499
1	.062500	2.5189	.025229	.075223 .0015953 .00031249
2	.083749	2.2954	.076875	.067457 .0036796 .00031249
3	.12500	2.0219	.15862	.030886 .015726 .00031249
old.I = .14180 .0057453 final.I = .69139 -.0055268				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.16386 .00078842 .00062499
1	.062500	2.5189	.025229	.075223 .0015953 .00031249
2	.083749	2.2954	.076875	.034833 .0014942 .00015624
3	.10937	2.1632	.11439	.032896 .0021329 .00015624
4	.12500	2.0219	.15862	.0000537 .0079676 .00015624
old.I = .067457 .0036796 final.I = .69139 -.0055268				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.16386 .00078842 .00062499
1	.062500	2.5189	.025229	.038546 .00068494 .00015624
2	.078125	2.4180	.047044	.036800 .00068812 .00015624
3	.083749	2.2954	.076875	.032896 .0021329 .00015624
old.I = .075223 .0015953 final.I = .75894 -.0019171				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.084137 .000052465 .00031249
1	.031250	2.6685	.0033577	.081022 .00044667 .00031249
2	.062500	2.5189	.025229	.036800 .00068812 .00015624
old.I = .16386 .00078842 final.I = .83433 -.00040523				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.084137 .000052465 .00031249
1	.031250	2.6685	.0033577	.041173 .00011247 .00015624
2	.048875	2.6036	.011036	.040020 .00028334 .00015624
3	.062500	2.5189	.025229	.032896 .0021329 .00015624
old.I = .081022 .00044667 final.I = .83433 -.00040523				

N	X	F	I	E
0	0.0000	2.7182	0.0000	.042371 .000003331 .00015624
1	.015625	2.7052	.00042642	.041986 .0000298 .000 5624
2	.031250	2.6685	.0033577	.040020 .00028334 .00015624
old.I = .084137 .000052465 final.I = .91558 -.000026373				
.99999 .00000001 ok				

Note:
answer = 1

§2 Fitting functions to data

One of the most important applications of numerical analysis is the representation of numerical data in functional form. This includes fitting, smoothing, filtering, interpolating, *etc.*

A typical example is the problem of table lookup: a program requires values of some mathematical function – $\sin(x)$, say – for arbitrary values of x . The function is moderately or extremely time-consuming to compute directly. According to the Intel tim-

ings for the 80x87 chip, this operation should take about 8 times longer than a floating point multiply. In some real-time applications this may be too slow.

There are several ways to speed up the computation of a function. They are all based on compact representations of the function – either in tabular form or as coefficients of functions that are faster to evaluate. For example, we might represent $\sin(x)$ by a simple polynomial¹⁸

$$\sin(x) \approx x (0.994108 - 0.147202x) , \quad (22)$$

accurate to better than 1% over the range $-\frac{\pi}{2} \leq x \leq \frac{\pi}{2}$, this requires but 3 multiplications and an addition to evaluate. This would be twice as fast as calculating $\sin(x)$ on the 80x87 chip¹⁹.

To achieve substantially greater speed requires table look-up. To locate data in an ordered table, we might employ binary search: that is, look at the x -value halfway down the table and see if the desired value is greater or less than that. On the average, $\log_2(N)$ comparisons are required, where N is the length of the table. For a table with 1% precision, we might need 128 entries *i.e.* seven comparisons.

Binary search is unacceptably slow – is there a faster method? In fact, assuming an ordered table of equally-spaced abscissae the fastest way to locate the desired x -value is **hashing**, a method for *computing* the address rather than finding it using comparisons. Suppose, as before, we need 1% accuracy, *i.e.* a 128-point table with x in the range $[0, \pi/2]$. To look up a value, we multiply x by $256/\pi \approx 81.5$, truncate to an integer and quadruple it to get a (4-byte) floating point address. These operations – including fetch to the 87stack – take about 1.5-2 fp multiply times, hence the speedup is 4-fold.

The speedup factor does not seem like much, especially for a function such as $\sin(x)$ that is built into the fast co-processor. However, if we were speaking of a function that is considerably slower to evaluate (for example one requiring evaluation of an integral or solution of a differential equation) hashed table lookup with interpolation can be several orders of magnitude faster than direct evaluation.

We now consider how to represent data by mathematical functions. This can be useful in several contexts:

- The theoretical form of the function, but with unknown parameters, may be known. One might like to determine the parameters from the data. For example, one might have a lot of data on pendulums: their periods, masses, dimensions, *etc.* The period of a pendulum is given, theoretically, by

$$\tau = \left(\frac{2\pi L}{g} \right)^{1/2} f \left(\frac{L}{r}, \frac{m_{\text{bob}}}{m_{\text{string}}}, \dots \right) \quad (23)$$

where L is the length of the string, g the acceleration of gravity, and f is some function of ratios of typical lengths, masses and other factors in the problem. In order to determine g accurately, one generally fits a function of all the measured factors, and tries to minimize its deviation from the measured periods. That is, one might try

$$\tau_n = \left(\frac{2\pi L_n}{g} \right)^{1/2} \left[1 + \alpha \frac{r_n}{L_n} + \beta \left(\frac{m_{\text{bob}}}{m_{\text{string}}} \right)_n + \dots \right] \quad (24)$$

for the n 'th set of observations, with g, α, β, \dots the unknown parameters to be determined.

- Sometimes one knows that a phenomenon is basically smoothly varying; so that the wiggles and deviations in observations are noise or otherwise uninteresting. How can we filter out the noise without losing the significant part of the data? Several methods have been developed for this purpose, based on the same principle: the data are represented as a sum of functions from a complete set of functions, with unknown coefficients. That is, if $\varphi_m(x)$ are the functions, we say (y_n are the data)

$$y_n = \sum_{m=0}^{\infty} c_m \varphi_m(x_n) \quad (25)$$

Such representations are theoretically possible under general conditions. Then to filter we keep only a finite sum, retaining the first N (usually simplest and smoothest) functions from the set. An example of a complete set is **monomials**, $\varphi_m(x) = x^m$. Another is **sinusoidal (trigonometric) functions**,

$$\sin(2\pi mx), \cos(2\pi mx), 0 \leq x \leq 1,$$

used in Fourier-series representation. **Gram polynomials**, discussed below, comprise a third useful complete set.

The representation in Eq. 25 is called **linear** because the unknown coefficients c_m appear to their first power. Thus, if all the data were to double, we see immediately that the c_m 's would have to be multiplied by the same factor, 2. Sometimes, as in the example of the measurement of g above, the unknown parameters appear in more complicated fashion. The problem of fitting with these more general functional forms is called **nonlinear** for obvious reasons. The **simplex algorithm** of Ch. 8 §2.3 below is an example of a nonlinear fitting procedure.

We are now going to write programs to fit both linear and nonlinear functions to data. The first and conceptually simplest of these is the **Fourier transform**, namely representing a function as a sum of sines and cosines.

§§1 Fast Fourier transform

What is a Fourier transform? Suppose we have a function that is **periodic** on the interval $0 \leq x \leq 2\pi$:

$$f(x + 2\pi) = f(x);$$

Then under fairly general conditions the function can be expressed in the form

$$f(x) = a_0 + \sum_{n=1}^{\infty} (a_n \cos(nx) + b_n \sin(nx)) \quad (26)$$

Another way to write Eq. 26 is

$$f(x) = \sum_{-\infty}^{+\infty} c_n e^{inx}. \quad (27)$$

In either way of writing, the c_n are called **Fourier coefficients** of the function $f(x)$. Looking, e.g., at Eq. 27, we see that the **orthogonality** of the sinusoidal functions leads to the expression

$$c_n = \frac{1}{2\pi} \int_0^{2\pi} f(x) e^{-inx} dx. \quad (28)$$

Evaluating Eq. 28 numerically requires – for given n – at least $2n$ points²⁰. Naively, for each $n = 0$ to $N-1$ we have to do a sum

$$c_n \approx \sum_{k=1}^{2N} f_k e^{-2\pi i n k / N}$$

which means carrying out $2N^2$ complex multiplications.

The **fast Fourier transform (FFT)** was discovered by Runge and König, rediscovered by Danielson and Lanczos and ~~re~~-rediscovered by Cooley and Tukey²¹. The FFT algorithm can be expressed as three steps:

- Discretize the interval, i.e. evaluate $f(x)$ only for

$$x_k = 2\pi \frac{k}{N}, \quad 0 \leq k \leq N-1.$$

Call $f(x_k) \equiv f_k$.

- Express the Fourier coefficients as

$$c_n = \sum_{k=0}^{N-1} f_k e^{-2\pi i n k / N}. \quad (29)$$

- With $w_n = e^{-2\pi i n / N}$, Eq. 29 is an $N-1$ 'st degree polynomial in w_n . We evaluate the polynomial using a fast algorithm.

20. to prevent aliasing.

21. See, e.g., D.E. Knuth, *The Art of Computer Programming* v. 2 (Addison-Wesley Publishing Co., Reading, MA, 1981) p. 642.

To evaluate rapidly the polynomial

$$c_n = P_N(w_n) \equiv \sum_{k=0}^{N-1} f_k(w_n)^k$$

we divide it into two polynomials of order $N/2$, dividing each of those in two, *etc.* This procedure is efficient only for $N = 2^\nu$, with ν an integer, so this is the case we attack.

How does dividing a polynomial in two help us? If we segregate the odd from the even powers, we have, symbolically,

$$P_N(w) = E_{N/2}(w^2) + w O_{N/2}(w^2). \quad (30)$$

Suppose the time to evaluate $P_N(w)$ is T_N . Then, clearly,

$$T_N = \lambda + 2T_{N/2} \quad (31)$$

where λ is the time to segregate the coefficients into odd and even, plus the time for 2 multiplications and a division. The solution of Eq. 31 is $\lambda(N-1)$. That is, it takes $O(N)$ time to evaluate a polynomial.

However, the discreteness of the Fourier transform helps us here. The reason is this: to evaluate the transform, we have to evaluate $P_N(w_n)$ for N values of w_n . But w_n^2 takes on only $N/2$ values as n takes on N values. Thus to evaluate the Fourier transform for all N values of n , we can evaluate the two polynomials of order $N/2$ for half as many points.

Suppose we evaluated the polynomials the old-fashioned way: it would take $2(N/2) \equiv N$ multiplications to do both, but we need do this only $N/2$ times, and N more (to combine them) so we have $N^2/2 + N$ rather than N^2 . We have gained a factor 2. Obviously it pays to repeat the procedure, dividing each of the sub-polynomials in two again, until only monomials are left.

Symbolically, the number of multiplications needed to evaluate a polynomial for N (discrete) values of w is

$$\tau_N = N\lambda + 2\tau_{N/2} \quad (32)$$

whose solution is

$$\tau_N = \lambda N \log_2(N). \quad (33)$$

Although the FFT algorithm can be programmed recursively, it almost never is. To see why, imagine how the coefficients would be re-shuffled by Eq. 30: we work out the case for 16 coefficients, exhibiting them in Table 8-1 below, writing only the indices:

Table 8-1 Bit-reversal for re-ordering discrete data

Start	Step 1	Step 2	Step 3	Bin ₀	Bin ₃
0	0	0	0	0000	0000
1	2	4	8	0001	1000
2	4	8	4	0010	0100
3	6	12	12	0011	1100
4	8	2	2	0100	0010
5	10	6	10	0101	1010
6	12	10	6	0110	0110
7	14	14	14	0111	1110
8	1	1	1	1000	0001
9	3	5	9	1001	1001
10	5	9	5	1010	0101
11	7	13	13	1011	1101
12	9	3	3	1100	0011
13	11	7	11	1101	1011
14	13	11	7	1110	0111
15	15	15	15	1111	1111

The crucial columns are “Start” and “Step 3”. Unfortunately, they are written in decimal notation, which conceals a fact that becomes glaringly obvious in binary notation. So we re-write them in binary in the columns Bin₀ and Bin₃ – and see that the final order can be obtained from the initial order simply by reversing the order of the bits, from left to right!

A standard FORTRAN program for complex FFT is shown below. We shall simply translate the FORTRAN into FORTH as expeditiously as possible, using some of FORTH’s simplifications.

One such improvement is a word to reverse the bits in a given integer. Note how clumsily this was done in the FORTRAN

```

SUBROUTINE FOUR1(DATA, NN, ISIGN)
C
C   from Press, et al., Numerical Recipes, ibid., p. 394.
C
C   ISIGN DETERMINES WHETHER THE FFT
C   IS FORWARD OR BACKWARD
C
C   DATA IS THE (COMPLEX) ARRAY OF DISCRETE INPUT
C   COMPLEX W, WP, TEMP, DATA(N)
REAL*8 THETA
J=0
DO 11 I=0,N-1      \ begin bit reversal
  IF (J.GT.I) THEN
    TEMP=DATA(J)
    DATA(J)=DATA(I)
    DATA(I)=TEMP
  ENDIF
  M=N/2
1  IF ((M.GE.1).AND.(J.GT.M)) THEN
    J=J-M
    M=M/2
    GO TO 1
  ENDIF
  J=J+M
11 CONTINUE      \ end bit reversal

MMAX=1      \ begin Danielson-Lanczos section
2  IF (N.GT.MMAX) THEN
    ISTEP=2*MMAX      \ executed lg(N) times
                        \ init trig recurrence

    THETA=3.1415926535897930/(ISIGN*MMAX)
    WP=CEXP(THETA)
    W=DCMPLX(1.00,0.00)
    DO 13 M=1,MMAX,2      \ outer loop
      DO 12 I=M,N,ISTEP      \ inner loop
        J=I+MMAX      \ total = N times
        TEMP=DATA(J)*W
        DATA(J)=DATA(I)-TEMP
        DATA(I)=DATA(I)+TEMP
      12 CONTINUE      \ end inner loop
    C
    W=W*WP      \ trig recurrence
    C
    13 CONTINUE      \ end outer loop
    MMAX=ISTEP
    GO TO 2
  ENDIF      \ end Danielson-Lanczos section
RETURN

```

program. Since practically every microprocessor permits right-shifting a register one bit at a time and feeding the overflow into another register from the right, **B.R** can be programmed easily in machine code for speed. Our fast bit-reversal procedure **B.R** may be represented pictorially as in Fig. 8-5 below.




Fig. 8-5 Pictorial representation of bit-reversal

Bit-reversal can be accomplished in high-level FORTH via

```

: B.R      ( n - - n' )    \ reverse order of bits
  0 SWAP   ( - - 0 n )    \ set up stack
  N.BITS 0 DO
    DUP 1 AND              \ pick out 1's bit
    ROT 2* +               \ left shift 1, add 1's bit
    SWAP 2/                \ right-shift n
  LOOP DROP ;

```

Note: **N.BITS** is a **VAR**, previously set to $v = \log_2(N)$

We will use **B.R** to re-order the actual data array (even though this is slightly more time-consuming than setting up a list of scrambled pointers, leaving the data alone). We forego indirection for two reasons: first, we have to divide by N (N steps) when inverse-transforming, so we might as well combine this with bit-reversal; second, there are N steps in rearranging and dividing by N the input vector, whereas the FFT itself takes $N \log_2(N)$ steps, *i.e.* the execution time for the preliminary N steps is unimportant.

Now, how do we go about evaluating the sub-polynomials to get the answer? First, let us write the polynomials (for our case $N = 16$) corresponding to taking the (bit-reversed) addresses off the stack in succession, as in Fig. 8-6 below.

$$\left. \begin{array}{l}
 w^7(f_7 + w^8 f_{15}) \\
 w^3(f_3 + w^8 f_{11}) \\
 w^5(f_5 + w^8 f_{13}) \\
 w^1(f_1 + w^8 f_9) \\
 w^8(f_8 + w^8 f_{14}) \\
 w^2(f_2 + w^8 f_{10}) \\
 w^4(f_4 + w^8 f_{12}) \\
 w^0(f_0 + w^8 f_6)
 \end{array} \right\} \left. \begin{array}{l}
 w^3(a_3 + w^4 a_7) \\
 w^1(a_1 + w^4 a_5) \\
 w^2(a_2 + w^4 a_6) \\
 w^0(a_0 + w^4 a_4)
 \end{array} \right\} \left. \begin{array}{l}
 w^1(b_1 + w^2 b_3) \\
 \\
 \\
 w^0(b_0 + w^2 b_2)
 \end{array} \right\} C_0 + w C_1$$

Fig. 8-6 The order of evaluating a 16 pt. FFT

We see that w_n^8 (for $N=16$) has only two possible values, ± 1 . Thus we must evaluate not 16×8 terms like $f_i + w^8 f_{i+8}$, but only 2×8 . Similarly, we do not need to evaluate 16×4 terms of form $f_i + w^4 f_{i+4}$, but only 4×4 , since there are only 4 possible values of w_n^4 . Thus the total number of multiplications is

$$2 \times 8 + 4 \times 4 + 8 \times 2 + 16 \times 1 = 64 \equiv 16 \log_2 16,$$

as advertised. This is far fewer than $16 \times 16 = 256$, and the ratio improves with N – for example a 1024 point FFT is 100 times faster than a slow FT.

We list the FFT program on page 191 below. Since **}FFT** transforms a one-dimensional array we retain the curly braces notation introduced in Ch. 5. We want to say something like

V{ n.pts FORWARD }FFT

where **V{** is the name of the (complex) array to be transformed, **n.pts** (a power of 2) is the size of the array, and the flag-setting words **FORWARD** or **INVERSE** determine whether we are taking a FFT or inverting one.

Now we test the program. Table 8-2 on page 192 contains the weekly stock prices of IBM stock, for the year 1983 (the 52 values have been made complex numbers by adding $0i$, and the table padded out to 64 entries (the nearest power of 2) with complex zeros)²². The first two entries (2, 64) are the type and length of the file. (The file reads from left to right.)

We FFT Table 8-2 using the phrase **IBM{ 64 DIRECT }FFT**. The power spectrum of the resulting FFT (Table 8-3) is shown in Fig. 8-7 on page 192 below.

22. This example is taken from the article "FORTH and the Fast Fourier Transform" by Joe Barnhart, *Dr. Dobb's Journal*, September 1984, p. 34.

```

\ Complex Fast Fourier Transform
\ Usage: Vector.name( N FORWARD ( INVERSE ) ) FFT
TASK FFT
\ check for presence of these extensions and load
FIND C+      0 = ?( FLOAD COMPLEX )
FIND 1ARRAY  0 = ?( FLOAD MATRIX.HSF )
FIND FILL    0 = ?( FLOAD FILEIO.FTH )
FIND TRIG     0 = ?( FLOAD TRIG )
\ if not there
DECIMAL
\ -----
\ auxiliary words
CODE SHR BX 1 SHR END-CODE
: LG2 ( n -- lg2(n) ) 0 SWAP (-- 0 n) SHR
  BEGIN ?DUP 0 >
  WHILE SHR SWAP 1+ SWAP REPEAT ;

\ VAR DIRECTION?
: FORWARD 0 IS DIRECTION? ;
: INVERSE -1 IS DIRECTION? ;

0 VAR N.BITS \ some VARs
0 VAR N
0 VAR MMAX
0 VAR I{

: C/N NS > F C/F ;
: NORMALIZE DIRECTION?
  IF C/N CPSWAP C/N CPSWAP THEN ;
\ end auxiliary words
\ -----
\ easy bit-reversal routines!
\ VAR LR
: BLR ( n -- n ) \ reverses order of bits
  0 SWAP (-- 0 n) \ set up stack
  N.BITS 0 DO DUP 1 AND \ pick out 1's bit
    ROT 2* + \ double sum and add 1's bit
    SWAP 2/ \ n - n/2
  LOOP DROP ;

: BIT.REVERSE 0 IS LR
  NO DO 1 BLR IS LR
  LR 1 < NOT ( LR > = 1? )
  IF I{ LR } G@L I{ 1 } G@L NORMALIZE
    I{ LR } GIL I{ 1 } GIL THEN
  LOOP ;
\ end bit-reversal (N times)

```

```

\ -----
\ main algorithm
CODE C+ 2 FLD. 1 FXCH. 3 FSUBR.
  1 FADDP. 1 FXCH. 3 FLD.
  1 FXCH. 4 FSUBR. 1 FADDP. 1 FXCH. END-CODE
  ( 87: w z -- w-z w+z )

: THETA F=PI MMAX S>F F/ DIRECTION? FSIGN ;

CREATE WP 16 ALLOT ONLY
: INIT.TRIG FINIT THETA EXP(*PI*) WP DCP1 C=1 ;

: NEW.W ( 87: w -- w' ) WP DCP@ C* ;

0 VAR ISTEP

: DO.INNER.LOOP
  DO MMAX 1 + IS LR
    CPDUP I{ LR } G@L C* I{ 1 } G@L
    CPSWAP C+ I{ 1 } GIL I{ LR } GIL
    ISTEP +LOOP ;

: }FFT ( adr n -- ) IS N IS I{
  1 IS MMAX
  N LG2 IS N.BITS
  FINIT
  BIT.REVERSE
  BEGIN
    N MMAX >
  WHILE
    INIT.TRIG MMAX 2* IS ISTEP
    MMAX 0 DO
      N 1 DO.INNER.LOOP NEW.W
    LOOP
    ISTEP IS MMAX
  REPEAT CPDPOP ;

: POWER 0 DO I{ 1 } G@L CABS CR 1. F. LOOP ;
\ power spectrum of FFT \ end of fft code
\ -----
\ an example
64 LONG COMPLEX 1ARRAY A{
: INIT.A A{ $ IBM.EX OPEN-INPUT FILL CLOSE-INPUT ;
INIT.A
A{ 64 DIRECT } FFT
\ end of example
\ -----

```

How do we know the FFT program actually worked? The simplest method is to inverse-transform the transform, and compare with the input file. The FFT and inverse FFT are given, respectively, in Tables 8-3 and 8-4 on page 193 below. Within roundoff error, Table 8-4 agrees with Table 8-2 on page 192.

Table 8-2 Weekly IBM common stock prices, 1983

2 64			
96.63 0.0	99.13 0.0	94.63 0.0	97.38 0.0
97.38 0.0	96.38 0.0	96.63 0.0	100.38 0.0
102.25 0.0	100.75 0.0	99.88 0.0	102.13 0.0
101.63 0.0	103.88 0.0	110.13 0.0	117.25 0.0
117.00 0.0	117.63 0.0	116.50 0.0	110.63 0.0
113.00 0.0	114.00 0.0	114.25 0.0	121.13 0.0
123.00 0.0	121.00 0.0	121.50 0.0	120.13 0.0
124.38 0.0	120.38 0.0	119.75 0.0	118.50 0.0
122.50 0.0	117.83 0.0	119.75 0.0	122.25 0.0
123.13 0.0	126.63 0.0	126.88 0.0	132.25 0.0
131.75 0.0	127.00 0.0	128.00 0.0	122.25 0.0
126.88 0.0	123.50 0.0	121.00 0.0	117.88 0.0
122.25 0.0	120.88 0.0	123.63 0.0	122.00 0.0
0.0 0.0	0.0 0.0	0.0 0.0	0.0 0.0
0.0 0.0	0.0 0.0	0.0 0.0	0.0 0.0
0.0 0.0	0.0 0.0	0.0 0.0	0.0 0.0




Fig. 8-7 Power spectrum of FFT of 1983 IBM prices (from Table 8-3)

Table 8-3 FFT of IBM weekly stock prices, 1983

0 5889.4889 0.0000000	32 3.1801882 0.0000000	16 7.2800000 -15.170043	48 7.2800000 15.170043
1 -1358.9239 582.94880	33 13.480087 44.888729	17 -27.548482 81.081904	49 89.775169 29.838997
2 -374.28417 988.84282	34 89.179887 32.348035	18 88.084281 117.80709	50 148.10401 -84.187837
3 305.86314 634.45819	35 99.788884 -25.042837	19 119.02930 71.422482	51 84.502433 -185.21282
4 327.40542 177.80118	36 39.788885 -83.410888	20 87.888384 -24.258874	52 -49.510889 -147.46129
5 125.21042 -8.0083304	37 -30.828173 -14.841442	21 1.8487881 -37.884341	53 -49.475231 5.4212083
6 -178.28804 39.830813	38 19.293210 29.322504	22 -11.327788 52.313049	54 80.279541 33.341928
7 -90.525482 228.27159	39 88.788382 22.908825	23 44.172229 110.77391	55 191.18481 -128.15400
8 150.24284 280.91506	40 107.71731 -34.415077	24 107.71731 34.415077	56 150.24284 -280.91506
9 191.18481 128.15400	41 44.172229 -110.77391	25 88.788382 -22.908825	57 -90.525482 -228.27159
10 80.279541 -33.341928	42 -11.327788 -52.313049	26 19.293210 -29.322504	58 -178.28804 -39.830813
11 -49.475231 -5.4212083	43 1.8487881 37.884341	27 -30.828173 14.841442	59 125.21042 8.0083304
12 -49.510889 147.46129	44 87.888384 24.258874	28 39.788885 83.410888	60 327.40542 -177.80118
13 84.502433 185.21282	45 119.02930 -71.422482	29 99.788884 25.042837	61 305.86314 -634.45819
14 148.10401 84.187837	46 68.084251 -117.80709	30 89.179887 -32.348035	62 -374.28417 -988.84282
15 89.775169 -29.838997	47 -27.548482 -81.081904	31 13.480087 -44.888729	63 -1358.9239 -582.94880

Table 8-4 Reconstructed IBM prices (inverse FFT)

0 96.630 0.0000	32 122.50 0.0000	16 117.00 -.000000000	48 122.25 .000000000
1 99.129 .000000705	33 117.82 -.000000245	17 117.82 -.000000518	49 120.87 .000000132
2 94.829 .000000023	34 119.75 .000000148	18 116.50 -.000000119	50 123.83 -.000000052
3 97.379 .000000227	35 122.24 -.000000379	19 110.82 .0000001375	51 121.99 -.000000944
4 97.380 .000000385	36 123.13 -.000000385	20 113.00 .0000001078	52 .000012884 -.000001078
5 96.379 -.000000362	37 126.82 .000000798	21 114.00 .000000975	53 -.000001277 -.000001434
6 98.630 .000001697	38 128.88 -.000000851	22 114.25 .000000378	54 -.000002880 -.000001224
7 100.38 .000002158	39 132.25 -.000001681	23 121.12 .000000317	55 -.000005683 -.000000915
8 102.25 -.000000000	40 131.75 -.000000000	24 123.00 .000000000	56 -.000001602 -.000000000
9 100.75 .000000799	41 127.00 -.000000373	25 121.00 .0000001130	57 -.000002985 -.000001630
10 99.890 -.000000271	42 128.00 .000000401	26 121.50 -.000001469	58 -.000010348 .000001339
11 102.12 -.000000078	43 122.24 -.000001011	27 120.12 .0000002522	59 -.000005828 -.000001711
12 101.82 .000000316	44 128.87 -.000000316	28 124.37 .000000027	60 -.000003018 -.000000027
13 103.88 -.000000043	45 123.50 .000000380	29 120.38 .000001281	61 -.000002253 -.000001583
14 110.13 -.000001815	46 121.00 .000001937	30 119.75 -.000000618	62 -.000003121 .000000295
15 117.25 -.000002237	47 117.87 .000001164	31 118.49 .000000573	63 -.000003713 .000000500

§§2 Gram polynomials

Gram polynomials are useful in fitting data by the linear least-squares method. The usual method is based on the following question: What is the "best" polynomial,

$$P_N(x) = \sum_{n=0}^N \gamma_n x^n, \quad (34)$$

(of order N) that I can use to fit some set of M pairs of data points,

$$\begin{Bmatrix} x_k \\ f_k \end{Bmatrix}, \quad k=0, 1, \dots, M-1$$

(with $M > N$) where $f(x)$ is measured at M distinct values of the independent variable x ?

The usual answer, found by Gauss, is to minimize the **squares of the deviations** (at the points x_k) of the fitting function $P_N(x)$ from the data — possibly weighted by the uncertainties of the data. That is, we want to minimize the **statistic**

$$\chi^2 = \sum_{k=0}^{M-1} \left(f_k - \sum_{n=0}^N \gamma_n x_k^n \right)^2 \frac{1}{\sigma_k^2} \quad (35)$$

with respect to the $N+1$ parameters γ_n .

From the differential calculus we know that a function's first derivative vanishes at a minimum, hence we differentiate χ^2 with respect to each γ_n independently, and set the results equal to zero. This yields $N+1$ linear equations in $N+1$ unknowns:

$$\sum_m A_{nm} \gamma_m = \beta_n, \quad n=0, 1, \dots, N \quad (36)$$

where (the symbol $\hat{=}$ means “is defined by”)

$$A_{nm} \hat{=} \sum_{k=0}^{M-1} (x_k)^{n+m} \frac{1}{\sigma_k^2} \quad (37a)$$

and

$$\beta_n \hat{=} \sum_{k=0}^{M-1} x_k^n f_k \frac{1}{\sigma_k^2} \quad (37b)$$

In Chapter 9 we develop methods for solving linear equations. Unfortunately, they cannot be applied to Eq. 36 for $N \geq 9$ because the matrix A_{nm} approximates a Hilbert matrix,

$$H_{nm} = \frac{\text{const.}}{n+m+1},$$

a particularly virulent example of an exponentially ill-conditioned matrix. That is, the roundoff error in solving 36 grows exponentially with N , and is generally unacceptable. We can avoid roundoff problems by expanding in polynomials rather than monomials:

$$\chi^2 = \sum_{k=0}^{M-1} \left(f_k - \sum_{n=0}^N \gamma_n p_n(x_k) \right)^2 \frac{1}{\sigma_k^2} . \quad (38)$$

The matrix then becomes

$$A_{nm} = \sum_{k=0}^{M-1} p_n(x_k) p_m(x_k) \frac{1}{\sigma_k^2} \quad (39a)$$

and the inhomogeneous term is now

$$\beta_n = \sum_{k=0}^{M-1} p_n(x_k) f_k \frac{1}{\sigma_k^2} \quad (39b)$$

Is there any choice of the polynomials $p_n(x)$ that will eliminate roundoff? The best kinds of linear equations are those with nearly diagonal matrices. We note the sum in Eq. 39a is nearly an integral, if M is large. If we choose the polynomials so they are orthogonal with respect to the weight function

$$w(x) = \frac{1}{\sigma_k^2} \theta(x_k - x) \theta(x - x_{k-1}) ,$$

where

$$\theta(x) = \begin{cases} 0, & x < 0 \\ 1, & x \geq 0 \end{cases}$$

then A_{nm} will be nearly diagonal, and well-conditioned.